

Narendra Devarasetty

Staff Data Engineer · Data Platforms · AI for Data

Naren9067@gmail.com · LinkedIn · GitHub · Tracy, CA 978-674-9067

SUMMARY

Seasoned Data Engineer with 15+ years of experience designing and re-architecting large-scale data platforms across product, finance, and ML. Proven track record of owning business-critical data systems end-to-end, from modeling and pipelines to reliability and cost efficiency, at companies operating at massive scale. Known for translating ambiguous business questions into durable, production-grade data systems that enable fast, confident decision-making.

EXPERIENCE

Airbnb · BizTech Data Engineering

Mar 2025 – Present

Staff Data Engineer · San Francisco, CA

- **Re-architected Unit Economics platform**, consolidating fragmented revenue, cost, and contribution-margin logic into canonical datasets powering finance close, forecasting, and product analytics across **10+ orgs**.
- **Delivered GPU efficiency initiative in 4 weeks**, partnering with Cloud Infrastructure to analyze **\$1.13M/month GPU spend** and identify **35–50% cost reduction potential (\$400K–\$600K/month)** across 3,000+ workloads.
- Served as **technical owner for cross-functional data alignment**, translating ambiguous business questions into durable, production-grade data models adopted by Finance, Product, and ML teams.

DoorDash · Core Data Engineering

May 2021 – Mar 2025

Staff Software Engineer, Data · San Francisco, CA

- **Owned backend architecture for AI-driven voice ordering system** delivered in **40 days**, including order-cart construction, LLM-optimized menu modeling, and inference integration for a high-visibility client demo.
- Built **synthetic data generation pipelines** enabling continuous fine-tuning of GPT-3.5 and LLaMA models, improving order accuracy and accelerating iteration without reliance on scarce real-world conversations.
- Built **LLM-powered data discovery and SQL generation tooling** used by **100+ analysts**, reducing ad-hoc analysis time by abstracting schema complexity through metadata and semantic search.
- Designed a **config-driven metrics platform** powering **1,000+ Tier-0 metrics** across dashboards, experimentation, and segmentation, eliminating bespoke SQL and standardizing metric logic.
- Built and scaled a **next-generation ETL platform** (Spark, Snowflake) with CI/CD, testing, observability, SLA enforcement, RBAC, and data quality checks, adopted by **20+ teams** and scaled to **2,500+ production workflows**.

- Led **Data Reliability Engineering (DRE)** sustaining **2+ years of SLA compliance** for 150+ Tier-0 metrics across 400+ upstream tables, reducing critical dashboard downtime by 3–4 hours.
- **Owned quarterly earnings metrics** and partnered with Finance, Product, and PwC to support audit and certification requirements.
- Led **company-wide data engineering cost efficiency initiative**, delivering **\$2M in savings in H2 2024** and defining a 2-year roadmap for sustainable cost governance.

Dropbox · Central Data Engineering

Jun 2019 – May 2021

Lead Data Engineer · San Francisco, CA

- Built **User & Team 360 canonical datasets** with **400+ attributes**, establishing a single source of truth supporting Product and Finance analytics across **700M+ users**.
- Created **Dropbox Usage DataMart** modeling file activity (adds, updates, shares, collaboration) over 1-day, 7-day, and 28-day rolling windows, ingesting and aggregating **billions of file events per day** at scale.
- Built **Finance DataMart for User and Team ARR**, integrating dimensional entities to deliver trusted financial reporting.
- Re-engineered **large-scale filesystem usage pipelines** using Spark, reducing ETL runtimes by **10+ hours**.
- Developed **reusable Airflow operators** for Hive/Spark workloads with built-in validation and structured logging, improving pipeline reliability and alerting.

Goldman Sachs · SecDiv Data Platform

Dec 2014 – May 2019

Data Engineer · New York, NY

- Built one of **Goldman Sachs' first Hadoop-based data lake platforms**, leveraging HDFS, Spark, Scala, and Pig to modernize large-scale analytics and batch processing.
- Designed **automated ETL frameworks** using Spark/Scala and Python, significantly improving data processing efficiency and operational reliability.
- Migrated **350+ legacy BI reports** to the modern data platform, optimizing batch job runtimes by 2+ hours and reducing maintenance overhead.
- Led **data onboarding initiatives to AWS and Snowflake**, enabling early cloud adoption and scalable analytics workflows within SecDiv.

Infosys · Fixed Income Data Platform

Jul 2007 – Nov 2014

Data Engineer · Nashua, NH

- Consolidated **500+ batch jobs into 25+**, improving reliability and maintainability of fixed-income data pipelines.
- Built real-time **Trade DataMart with 5-minute SLA**, implemented Oracle VPD for data security, and led early Hadoop POCs for scalable analytics.

EDUCATION

Anna University, India

May 2003 – Apr 2007

Bachelor of Engineering, Electronics and Communication